

TETHER: Tethered Causal Temporal Refinement for Frozen Surgical Stereo

Anonymous Authors

Abstract—Modern stereo estimators are applied independently to each frame and can fluctuate under specularities, weak texture, tissue deformation, tool occlusions and camera motion. Video stereo methods fix this by building temporal reasoning into the estimator, so a new backbone means a new temporal model. We ask the opposite question: can one causal temporal module be trained once and attached, unchanged, to frozen and heterogeneous stereo estimators?

We present TETHER, a causal plug-and-play refiner that answers it. One frozen 154,874-parameter checkpoint, trained on three estimators and one dataset, improves *every* estimator on *every* arena we evaluate: five architectures, two excluded from its training, over development, held-out and out-of-domain splits — all fifteen backbone-arena cells in mean EPE, with the sequence-level exceptions named where they occur. Zero-shot on DRENDS, which no part of the module has seen, disparity EPE, Bad3 and metric depth MAE and RMSE improve in all 25 recording-backbone cells, by -5.41 , -19.51 , -4.02 and -3.47% , nothing adapted or tuned for the domain; the two unseen estimators we cannot separate from the three seen ones, which is the claim.

The module reads only what a black box emits: preceding disparity motion-aligned with frozen SEA-RAFT, and a shared 38-channel evidence space of disparity, motion and image geometry — no learned appearance features, no frozen backbone, no future frame — feeding a reset-and-fusion head that predicts a bounded interpolation between the current estimate and that aligned memory. One of its inputs, a forward-backward flow-consistency cue, is needed in training but not at inference, so the module deploys at ~ 16 of its ~ 29 ms on the unchanged checkpoint. Under a matched state machine the gain is not temporal smoothing: 3.69% mean EPE against 0.59% for the best fixed or EMA blend, which buys its mean error with bad pixels while the learned gate does not. Accumulated into a map the refined geometry reduces world-frame disagreement in all 20 held-out cells and increases it on development; we identify the term that decides which. We additionally introduce a refinement-centric protocol scoring introduced error alongside recovered error, and apply it to our own claims.

I. INTRODUCTION

Metric three-dimensional perception underpins camera navigation, tool-tissue distance estimation and deformation analysis in robot-assisted surgery, and stereo endoscopy is attractive because depth comes from images the platform already provides. It is also hard: endoscopic scenes carry weak texture, non-Lambertian tissue, specularities, blood, smoke, deformable anatomy, rapid viewpoint change and tool occlusion [18], [19]. Modern estimators [11]–[15] give increasingly strong frame-wise disparity, but nothing in a frame-wise model requires its predictions to cohere over time. Video stereo methods build temporal reasoning into the model [1]–[4], at the price of a temporal architecture,

model-specific representations, future observations, or co-design with the estimator.

We study a different deployment setting: *can one causal temporal module be trained once and attached unchanged to heterogeneous frozen stereo estimators?* Backbone and temporal logic can then evolve independently, a new estimator replacing the previous one without the module needing its cost volume, hidden state, confidence head or training pipeline. We keep one inductive bias from the video-stereo literature — decide whether motion-aligned history should influence the current estimate, and constrain the correction to a convex interpolation between the two — and redesign everything around it for a black-box interface.

Two questions follow. Does the learned head add anything simple temporal averaging would not? Matched blends recover only a small fraction of its gain while an oracle shows the aligned memory carries far more than the policy exploits, so the core problem is not obtaining temporal evidence but deciding how much of it to use. And how should such an intervention be evaluated? Smoothness is not correctness, and lower mean EPE hides whether refinement introduces new bad pixels or rare catastrophic frames, so we score relative to the frozen prediction — and, because what a robot consumes is a map, carry the same intervention into an accumulated reconstruction and measure it there.

Our contributions are three.

(i) **One module, every backbone, every domain.** We measure a single frozen checkpoint across the full backbone $\times$ domain matrix — five estimators, two excluded from its training, over development, held-out and out-of-domain splits, including the joint unseen-backbone, unseen-domain cell usually left unmeasured — and it improves all fifteen cells, with the unseen estimators we cannot separate from the seen ones.

(ii) **A black-box interface that makes that possible.** We formulate surgical temporal stereo refinement as a *causal plug-and-play adaptation problem* over frozen, heterogeneous estimators, needing neither future frames nor backbone internals, and instantiate it as TETHER: a 154,874-parameter reset-and-fusion head over causal SEA-RAFT alignment, a shared 38-channel external evidence space, support-confined convex fusion and bounded corrected-state recurrence. Under a matched state machine the gain is not temporal smoothing.

(iii) **A protocol that scores the intervention, not the prediction.** We measure introduced error alongside recovered error under fixed, prediction-independent support, and it changes conclusions: on development, every spatially uniform blend recovering a meaningful share of the gain buys

it by *raising* the bad-pixel rate, while the one unlearned rule that varies in space does not. That trade is a property of the split as much as of the policy: out of domain, where the raw prediction is weak, none of the uniform blends pays in bad pixels either. Applied to our own claims the protocol is equally unkind: degenerate history-replay wins the no-reference metric, the head’s weight inverts its relation to harm across splits, and what a map fuses improves on one split while degrading on the other, for a reason we can name.

II. RELATED WORK

a) Deep stereo matching.: Stereo has moved from cost-volume regression to iterative recurrent estimation [11], [12] and broad-domain zero-shot models [13]–[15]. Our work is orthogonal: we neither modify the estimator nor assume an architecture.

b) Temporal stereo and depth.: CODD introduced the reset/fusion mechanism we build on [1]; DynamicStereo [2], PPMStereo [4] and TC-Stereo [6] follow, all co-designed with the estimator, so a new backbone means a new temporal model. Match Stereo Videos exposes BiDAStabilizer as a plug-in on a frozen estimator [3], but propagates bidirectionally. Monocular video depth is made consistent by diffusion [24] or stabiliser networks [16] — a different output space, none causal at clip level — while online weight adaptation [25] is the opposite of our constraint.

c) What can actually be compared.: We are aware of no published method that is simultaneously causal, black-box and estimator-agnostic. CODD [1], whose two-gate decision we build on, is causal but not black-box: its fusion network consumes matching features from its own stereo network, which a cached third-party backbone does not expose. The 142-channel configuration of Sec. V-J began as the closest reproduction our interface admits — its projection replaced by our warp, a frozen ResNet-18 standing in for the estimator features — so Table IV is the nearest thing to running CODD here. Among black-box candidates StereoDiffusion [7] ships no implementation and [8], [9] publish code whose weights are no longer retrievable (audited in the supplement). Two are runnable and we report both: BiDAStabilizer [3], which refines a frozen estimator as we do but bidirectionally, and TC-Stereo [6], whose rigid pose warp needs a privileged input our interface does not assume but SCARED-C supplies, so we run it as its authors intended (Sec. V-I).

d) Evaluating temporal refinement.: Stereo benchmarks report frame-wise geometry and video methods add consistency measures, but neither characterises an intervention: a prediction can grow smoother without growing more correct, and a lower mean error can hide newly created failures. Recovered and introduced error must both be measured, which connects to selective prediction [17] and to the uncertainty literature [26]. Our contribution is not a new metric but a common protocol under fixed, prediction-independent support.

e) Surgical stereo geometry.: SCARED [18] and DRENDS [22] supply the endoscopic stereo and time-of-flight geometry we evaluate on; D4D [21] enters only as the

substrate for a no-reference control; SERV-CT [19] not at all, its pairs being non-consecutive. Downstream reconstruction and tracking pipelines consume exactly the per-frame geometry we refine [20], [27], which is why we treat refinement as an intervention on that geometry and, in Sec. V-H, measure it after accumulation rather than only before.

III. METHOD

A. External Causal Interface

A frozen estimator S produces positive-left disparity d_t^{raw} and its validity M_t^{raw} . The temporal module never reads the internal representation of S ; it consumes only externally observable quantities — RGB, disparity, validity, optical flow and evidence derived from them. At time t the cache holds the preceding left image, the preceding raw disparity and validity, the previous fused state and its age; nothing from $t + 1$ onwards exists. The map to learn is $d_t^{\text{fused}} = T_\theta(d_t^{\text{raw}}, m_{t-1}, I_t^L, I_t^R, I_{t-1}^L)$.

B. Causal Motion Alignment

Frozen SEA-RAFT [5], [23] estimates both directions between the two already observed left images; the reverse direction is used only to measure reliability, and bidirectional *flow estimation* between two observed frames is not bidirectional temporal inference. Previous geometry is pulled onto the current grid, $\tilde{m}_t(x) = m_{t-1}(x + F_{t \rightarrow t-1}(x))$, by bilinear sampling with zero padding; warp support S_t records whether the source coordinate lies inside the preceding image, and historical validity is sampled on the same grid. Pulling the reverse flow the same way gives the forward-backward residual e_t^{FB} and the reliability cue $C_t^{FB} = S_t \exp(-e_t^{FB}/\rho_t)$ with ρ_t scaled by the flow magnitudes. The zero padding is not a detail: it makes $\tilde{m}_t$ exactly zero outside the source image, and Sec. V-K shows what happens when that value reaches the output.

C. Shared External Evidence

All estimators are mapped into one 38-channel tensor on the cache grid, built from raw disparity, aligned memory, the two images, the previous left image, the flow, C_t^{FB} , the warp support and the sampled historical validity. Twenty-five channels are disparity-geometry correlations formed at quarter resolution — 16 self-correlations of the two candidates and 9 raw-to-aligned cross-correlations over a dilated 3×3 neighbourhood — and 13 are formed on the grid directly: the two disparities, their signed and absolute residual, 6 motion components and 3 RGB. Normalisation is fixed before training. Nothing in this space is learned, no network evaluates it, and no ground truth, backbone identity or future frame enters it. An earlier configuration added 104 quarter-resolution channels derived from a frozen ResNet-18 [10] — 64 compressed feature maps and 40 confidence curves, supports and appearance correlations formed from them — for 142 in total. The channel ladder of Table IV reads off this decomposition: dropping only the 64 feature maps gives 78, dropping all 104 is TETHER’s 38, and dropping its 25 disparity correlations leaves the 13 direct




Fig. 1. One time step, and the evidence it reads. **Left:** dashed grey is frozen; the single orange block is the only trained component and emits one scalar per pixel, w_t , rather than geometry. **Right:** the 38 channels — 25 quarter-resolution disparity correlations, upsampled, and 13 native — none carrying ground truth, backbone identity or a future frame.

channels. Sec. V-J reports the 142-channel configuration as an ablation, because it costs 157,504 frozen parameters and three forward passes per frame and does not survive a three-seed comparison. Exact channel construction and constants are in the supplement.

D. Reset-and-Fusion Decision Head

The two-gate decomposition of the temporal decision, and the constraint that the correction be convex in the two candidates, are CODD’s [1]; the architecture, the motion pathway, the evidence space, the supervision and the training are not.

A cache-grid pathway and a quarter-resolution context pathway (Fig. 1) feed a full-resolution reset branch predicting $r_t = \sigma(z_t^r)$ and a coarse fusion branch predicting $f_t = \uparrow_4 [\sigma(z_t^f)]$, giving the effective temporal weight $w_t = r_t f_t$ from 154,874 trainable parameters. Running the context branch coarsely is a compute choice on that branch, $16\times$ cheaper than full resolution, which Sec. V-J reports as an ablation. Layer widths are in the supplement.

Temporal evidence is only defined where the aligned memory is usable, so the per-pixel *support* is the conjunction of current raw validity, aligned temporal validity and warp support,

$$S_t = M_t^{\text{raw}} \wedge \widetilde{M}_t \wedge \mathcal{S}_t, \quad (1)$$

and we confine the intervention to it. Final disparity is

$$d_t^{\text{fused}} = \begin{cases} (1 - w_t) d_t^{\text{raw}} + w_t \widetilde{m}_t, & S_t = 1, \\ d_t^{\text{raw}}, & S_t = 0. \end{cases} \quad (2)$$

Because $w_t \in [0, 1]$ the fused disparity lies between the two candidates on the support, and off it the module is the exact identity on raw geometry. Applying (1) to the output and not only computing it is load-bearing (Sec. V-K). Reset output bias is initialized to -2 and fusion to zero, so initial intervention is conservative.

E. Bounded Corrected-State Recurrence

After refinement $m_t = d_t^{\text{fused}}$, so alignment and fusion errors can become future evidence. We bound corrected-state lifetime: the first output of a sequence is raw, a re-anchor restores previous raw disparity as state, and otherwise the previous fused output is propagated.

The shipped policy permits four corrected-state uses before a raw re-anchor, so H denotes state lifetime rather than a temporal input window; at lifetime H the mean age of the state in use is $(H - 1)/2$ frames, the quantity Sec. V-H finds the accumulated cost tracks. The method remains one-step causal: $\{I_{t-1}, I_t, m_{t-1}\} \rightarrow d_t^{\text{fused}}$.

F. Training Objective

For target disparity g_t let $\Delta_t = |\widetilde{m}_t - g_t| - |d_t^{\text{raw}} - g_t|$, so $\Delta_t < 0$ marks pixels where aligned memory is the better candidate; training support is the intersection of GT coverage, current validity, aligned temporal validity and warp support. The unweighted objective $\mathcal{L} = \mathcal{L}_{\text{disp}} + \mathcal{L}_{\text{reset}} + \mathcal{L}_{\text{fusion}}$ pairs a Huber term on the fused disparity with two decision terms supervising each branch against the sign of Δ_t at native margins 5 and 1 px, scaled to the cache grid, plus a weak pull of the fusion weight towards 0.5 where the two candidates are within the fusion margin. No explicit temporal-consistency objective is used: temporal behaviour is a consequence of the decision, never a training target.

IV. EXPERIMENTAL PROTOCOL

A. Training

The temporal refiner is trained on temporally curated SCARED-C data derived from SCARED [18], using frozen predictions from S²M²-S [14], RAFT-Stereo [11] and Stereo Anywhere [13]; CREStereo [12] and Fast-FoundationStereo [15] are excluded from temporal-head training, and no backbone identity is ever given to TETHER.

Training uses D1, D3 and D6 with D2 for development and D7 held out, over 8,616 frames giving 8,606 causal pairs per backbone from three seen backbones, in non-overlapping $H = 4$ clips, batch size 4, 12 epochs of AdamW at 2×10^{-4} with weight decay 10^{-4} and gradient clipping 1.0, selecting on development fused EPE. The budget is not a truncation: selection lands at epoch 8, 11 and 11 across three seeds, all inside the schedule, and a locked-recipe run to 47 epochs oscillates without trend from epoch 13 with no later epoch improving on it. One ablation (A4) diverges after epoch 10, which best-validation selection catches. Details in the supplement.

B. External Evaluation

DRENDS [22] supplies dynamic robotic-endoscopy stereo with calibration and metric geometry; we report five curated

recordings, 7476 frames each, with all five estimators. Its disparity reference comes from metric ToF geometry through stereo calibration, so it is useful metric evidence but not independent stereo ground truth, and the ToF reference is itself temporally filtered. The zero-shot claim applies to the SCARED-trained checkpoint: no DRENDS data enters training, selection or calibration.

C. Unified Temporal-Refinement Evaluation

Evaluation support cannot depend on the prediction: $M_{\text{eval}} = M_{\text{GT}} \cap M_{\text{protocol}}$, and invalid predictions on valid support are penalised rather than silently removed. Temporal fidelity is change error DTCE on that support at lags $k \in \{1, 2, 4, 8\}$ ($k=4$ unless stated), scored against the *reference* change rather than stillness, so the degenerate replay of Sec. VII does not win it. Intervention harm and benefit come from $\delta e = e^{\text{fused}} - e^{\text{raw}}$: the masses $H^+ = \mathbb{E}[\max(\delta e, 0)]$ and $B^+ = \mathbb{E}[\max(-\delta e, 0)]$, their rates, newly broken and recovered bad pixels per threshold, and identity preservation. Aggregation is named at every number: macro is the mean of per-cell relative reductions, pooled sums before the ratio. The rest of the family — spatial and metric-depth metrics, per-frame tails, risk-coverage [17] — is in the supplement.

V. RESULTS

The claim is one module across backbones and domains, so the results are ordered that way: the transfer matrix, the seed evidence that it is stable, why it works, what it does to a map, and what it costs.

A. Cross-Backbone Development Behaviour

$H = 4$ improves all five estimators in mean D2 EPE (Table I), from 1.41% to 8.19% averaging per-sequence reductions — 1.70% to 8.04% from the pooled cells the table prints — and all six unseen-backbone sequence groups improve; on D7 the mean excluding Stereo Anywhere is 5.15%. One D2 cell worsens, S²M²-S on keyframe 2, by 0.28%: the first of three sequence-level exceptions in this paper, the others being Vid13 (Sec. V-I) and the development map (Sec. V-H).

B. Held-Out Cross-Backbone Confirmation

We next evaluate the already frozen $H = 4$ policy on held-out D7. All five backbones improve, including both unseen during temporal-head training: mean EPE decreases $0.4921 \rightarrow 0.4651$, or 5.49% pooling before the ratio and 5.40% averaging per-cell reductions.

The improvement is not restricted to the mean: Bad1, RMSE and P99 decrease for every backbone on both splits, so the EPE gain is not bought by a worse bad-pixel rate or a broader tail, though new bad pixels do appear and are measured below. InvalidRate stays zero in all ten SCARED-C backbone-split evaluations — a property of SCARED-C, not of the module (Sec. V-K).

C. External Zero-Shot Geometry on DRENDS

We evaluate the frozen checkpoint on five DRENDS recordings with all five estimators, two excluded from temporal-head training, 7476 frames each, nothing adapted or tuned for the domain. Disparity EPE, Bad3, metric depth MAE and RMSE improve in *every* one of the 25 recording-backbone cells, by -5.41 , -19.51 , -4.02 and -3.47% on average. Residual variance is set more by the recording than the estimator but not uniformly: the EPE change spans 0.87 points across backbones on Vid10 against 3.29 on Vid13.

Averaged over backbones the mean EPE reduction is 3.96% seen against 2.67% unseen on D2, 5.80 and 4.79% on D7, 5.57 and 5.16% on DRENDS: every combination positive. We can separate neither group from the other and report the margins rather than a test — with three seen and two unseen backbones and per-backbone margins from a single seed, absence of a detectable gap is not evidence of no gap. We read no ordering into the numbers either: on D2 the seen advantage is carried entirely by Stereo Anywhere and reverses without it, so we withdraw the claim we drew about which axis costs more. This establishes zero-shot geometric transfer under a ToF-derived reference; it does not establish universal robustness, and DRENDS remains a non-independent reference.

D. Seed Variance, and What It Reveals About Selection

We pre-registered a seed study before any result existed: only the seed varies, resampling is at sequence granularity rather than pixel, and selecting the best seed is prohibited.

Pooled over five backbones the relative Bad1 reduction is $12.73 \pm 0.32\%$ on D2 and $5.84 \pm 0.22\%$ on held-out D7; EPE gives $4.13 \pm 0.38\%$ and $5.61 \pm 0.42\%$, where $\pm$ is the standard deviation over three seeds throughout this paper. A paired sequence-level bootstrap is positive in every resample but the cells are not independent draws, so we report it as descriptive only (supplement); the distribution-free statement the splits can carry is that the effect is positive in 3/3 and 4/4 sequences, a sign test at $p = 0.125$ and $p = 0.0625$. No split here reaches $p < 0.05$ and none can at these sequence counts: we claim consistency of direction, never significance.

The spread is the strongest argument for the shipped configuration, and not what we found first. The 142-channel ablation is stable held out (5.15 ± 0.34 Bad1) and unstable on development, its seeds falling monotonically over a 5.03-point range. We read that as selection optimism, since checkpoints are selected on D2. Removing the learned evidence removes the instability instead — 0.32 points of D2 seed deviation against 2.53, with no change to the selection rule — so the instability was the evidence space and not the protocol.

E. Does the Learned Temporal Policy Matter?

We isolate the learned decision policy on D2, seed 0. All variants share the frozen predictions, the alignment, the support contract and the evaluation code, and all run unconfined so the comparison shares one implementation (confinement moves D2 EPE by 1.2×10^{-5} px). Figures are

TABLE I

ONE FROZEN TETHER CHECKPOINT, SEED 0, ACROSS FIVE ESTIMATORS AND THREE ARENAS; EVERY EPE AND BAD1 CELL IS RAW→TETHER, LOWER IS BETTER, AND THE LOWER-RIGHT BLOCK IS THE JOINT UNSEEN-BACKBONE, UNSEEN-DOMAIN SHIFT USUALLY LEFT UNMEASURED. DTCE IS THE MACRO REDUCTION IN CHANGE ERROR AT $k=4$, HIGHER IS BETTER; ON DEVELOPMENT IT IS BACKBONE-DEPENDENT (SEC. V-F). CELLS ARE POOLED PER BACKBONE; PER-CELL PROSE FIGURES AND DRENDS BAD3 RECONSTRUCT FROM THE SUPPLEMENT’S PER-SEQUENCE TABLES.

Backbone	Status	D2 development			D7 held out			DRENDS, out of domain		
		EPE	Bad1 (%)		EPE	Bad1 (%)	DTCE	EPE	Bad1 (%)	DTCE
S ² M ² -S	seen	.2704 → .2658	1.39 → 1.27		.5176 → .4992	6.44 → 6.03	6.5	1.3477 → 1.2812	62.55 → 58.98	11.3
RAFT-Stereo	seen	.2691 → .2623	1.60 → 1.43		.4671 → .4296	6.09 → 5.75	7.3	1.3125 → 1.2371	58.83 → 55.19	8.5
Stereo Anywhere	seen	.3022 → .2779	2.18 → 1.74		.4774 → .4408	6.23 → 5.77	7.9	1.3685 → 1.2834	62.73 → 57.99	13.9
CREStereo	unseen	.2668 → .2583	1.57 → 1.40		.5165 → .4912	6.35 → 6.05	6.4	1.3027 → 1.2296	55.17 → 52.21	11.8
Fast-FoundationStereo	unseen	.2725 → .2658	1.60 → 1.44		.4819 → .4645	5.92 → 5.65	4.3	1.2838 → 1.2228	54.88 → 52.06	9.6

TABLE II

MATCHED-POLICY CLOSURE ON D2, SEED 0, FIFTEEN BACKBONE-SEQUENCE CELLS POOLED BEFORE EACH RATIO; REDUCTIONS IN %, HIGHER IS BETTER, DTCE AT $k=4$. EVERY ROW SHARES THE PREDICTIONS, ALIGNMENT, SUPPORT CONTRACT, $H=4$ STATE MACHINE AND EVALUATION CODE — ONLY THE POLICY CHANGES; THE FIRST SIX LOAD NO CHECKPOINT, AND UNBOUNDED RECURRENCE IS OUR OWN HEAD WITH THE RE-ANCHOR REMOVED. TETHER’S ROW IS SEED 0, THE WEAKEST OF ITS THREE (3.69 AGAINST 4.13 ± 0.38), SO EVERY MARGIN HERE UNDERSTATES IT.

Policy	EPE	Bad1	DTCE _{$k=4$}	NewBad1
Fixed $w=0.5$ ($\equiv$ EMA2)	0.59	−2.64	6.26	0.214
EMA, three frames	0.30	−8.88	4.88	0.379
FB-confidence as the weight	−2.17	−32.35	−7.10	0.831
Warped raw previous, $H=1$	−0.03	−6.39	−4.33	0.373
Tuned residual rule	0.87	0.33	2.88	0.088
Aligned-memory propagation	−3.56	−41.70	−13.55	1.072
Unbounded corrected recurrence	−19.17	−143.48	−28.13	2.858
TETHER, $H=4$	3.69	12.84	4.99	0.063
Raw-vs-memory oracle (ceiling)	12.74	21.87	10.01	0.000

pooled over the fifteen backbone-sequence cells, whose raw EPE is 0.276202.

The fifteen policies (Table II, full grid in the supplement) are the horizon sweep and continuous recurrence for the learned head, fixed $w \in \{0.1, 0.2, 0.3, 0.5\}$ rising monotonically to 0.59%, EMA over two and three frames, forward-backward confidence as the weight, aligned-memory propagation, warped raw previous at $H=1$, and a tuned two-parameter rule on the raw-versus-memory residual swept over seventeen (α, τ) points. Per backbone, against whichever fixed or EMA policy is best for that backbone, the ratio ranges $3.9\times$ to $9.9\times$: the absolute gain depends on the estimator, the margin over smoothing does not.

Three rows carry more than the mean. Forward-backward confidence is the informative failure — one of TETHER’s own cues, yet used directly as the weight it worsens EPE by 2.17% and Bad1 by 32.35%. The uniform smoothing policies buy mean error with bad pixels, the more so the harder they blend: $w=0.3$ costs 0.18% more bad pixels for 0.49% of EPE, the best fixed blend 2.64% for 0.59%, EMA3 8.88% for 0.30%. The tuned residual rule limits that statement and we report it as such: it recovers roughly a quarter of our gain and does *not* pay in bad pixels. So the trade is not a property of blending but of blending *uniformly* — the moment a rule varies in space it stops buying mean error with tails. That

is the closest an unlearned policy comes, and where the distance shows: 0.87% against 3.69% of EPE, breaking more pixels doing less. What the learned head adds is not varying in space; it is knowing where.

The closure is development work, and a reviewer is right to ask whether the margin survives the split. We repeat it out of domain, on the arena the freeze permits: on DRENDS with RAFT-Stereo, the best unlearned policy recovers 0.63% of EPE against our 5.80% (RAFT-Stereo alone, pixel-pooled on the closure’s support; the five-backbone per-cell mean is 5.41) — a ratio of $9\times$ against $6\times$ on D2. The denominator barely moves between arenas: uniform smoothing recovers the same small amount wherever it runs, so the widening is entirely ours, and it tracks how much raw error there is to recover (0.276 px on D2 against 1.286 out of domain, both on the closure’s own support), the same quantity Sec. V-H finds decides the accumulated benefit. Where the raw prediction is this weak the uniform blends no longer pay in bad pixels, so that trade is stated on D2 alone.

F. Why $H = 4$, and Why Recurrence Must Be Bounded

Recurrence must be bounded: with no re-anchor the intervention inverts, EPE degrading by 19.17% and Bad1 by 143.48% (Table II). Within the bounded range EPE rejects $H=1$, by 1.24 points, and does not separate the other three, whose spread is 0.21 (3.69, 3.73, 3.52%), while newly broken pixels rise monotonically from 0.052% to 0.128%. We froze $H=4$ before D7 opened on that risk-gain trade, and it is the last horizon at which change error also improves at lags $k \in \{2, 4, 8\}$ (4.99% at $H=4$ against −3.64% at $H=6$ for $k=4$). Two caveats. At $k=1$ DTCE improves at every horizon we tried, so the sign change belongs to the longer lags, not the metric. And on D2 the improvement is carried by one backbone and is *negative* for three: +14.0% for Stereo Anywhere and +0.8% for CREStereo against −0.3, −2.0 and −2.3%. A criterion whose sign one of five backbones decides cannot carry an operating point, so the horizon rests on breakage; on both confirmation splits DTCE is positive for every backbone (Table I). Sec. V-H sweeps the same horizon against accumulated geometry and finds the same risk structure from an independent measurement.

G. Refinement as an Intervention

Across both splits and all five backbones TETHER recovers $3.8\text{--}11.1\times$ as many Bad1 pixels as it introduces,

B^+/H^+ between 1.39 and 3.25. Per frame, 56% of the 15,535 held-out frames improve against 14% worsened (45 against 18% on development, the rest exact Bad1 ties on cleaner frames), the worsened band hugging the diagonal while the improved mass extends far from it: the count-versus-magnitude asymmetry behind B^+/H^+ .

TABLE III

ACCUMULATED GEOMETRY ON RAFT-STEREO: CHANGE IN WORLD-FRAME DISAGREEMENT, AS % OF THE EXCESS OVER THE GROUND-TRUTH FLOOR, POSITIVE BETTER. ALL ROWS SHARE THE POSES, THE ACCUMULATION AND THE FIXED SUPPORT; THE GRID IS COMPLETE (3 D2 AND 4 D7 CELLS PER POLICY).

Policy	D2 development	D7 held out
Copy previous (stillness control)	−54.0	+2.8
Ungated warp, full weight	−37.6	+7.2
TETHER, $H=1$	−3.6	+8.9
TETHER, $H=2$	−8.6	+10.4
TETHER, $H=4$ (shipped)	−15.5	+11.7
TETHER, $H=6$	−20.2	+12.2

H. What a Map Fuses, Not What a Frame Scores

Everything above scores a frame. What a robot consumes is a map, so we measure that too: SCARED-C’s corrected poses let geometry be accumulated in one world frame and the disagreement between views of a point measured, reported as a percentage of the excess over the ground-truth floor. On held-out D7 refinement reduces it by 11.7%, and the sign is not carried by one estimator: across five backbones and four sequences all 20 cells improve, by +7.2 to +17.8%.

Our protocol requires any new measure to face the controls that defeat the no-reference score. Copy-previous, which removes 97.3% of that score, is here the *worst* policy in Table III: it degrades the map by 54.0% on D2 and returns 2.8% on D7 against our 11.7%. Disagreement between views is not reducible by repetition, because a replayed frame is wrong from the new viewpoint, so the measure survives the control our own contribution (iii) demands of temporal metrics.

The horizon sweep names the condition, on RAFT-Stereo across the complete grid; the five-backbone confirmation exists at the shipped horizon only (Sec. V-B), so the mechanism below is established on one estimator and its sign confirmed on five. On D2 the cost grows monotonically with state lifetime and decelerates — 3.6% at $H=1$, one warp and no accumulation, to 20.2% at $H=6$ — so a map pays not for the warp but for carrying it. The ungated warp at full weight costs 37.6% where the gated head at $H=1$ costs 3.6%, and on D7 already returns 7.2 of our 11.7: on a map most of the gain is alignment and most of the gate’s value is damage avoided, the same division Sec. VI finds per pixel. Since alignment residual is a property of the motion while raw excess differs twofold between the splits (0.171–0.290 mm against 0.518–1.113), whether that residual is bought back is decided by the split.

Both splits are monotone in H and neither turns, so the horizon is set by an exchange rate rather than a peak. Held-

out benefit saturates while development cost keeps growing — successive steps buy +1.5, +1.3 and +0.5 points on D7 for −5.0, −6.9 and −4.7 on D2 — so the price of lifetime rises from 3.3 to 9.4 points of development error per point of held-out gain, making $H=6$ the best cell of the D7 column and the worst purchase on the sweep. We name it because the D7 column read alone would suggest a longer horizon. $H=4$ was frozen on the per-frame trade before D7 opened and before this measurement existed, and the map is the fourth criterion to place it there: against $H=6$, frame EPE is unreadable at 0.21 points of spread, newly broken pixels rise, change error changes sign, and the map costs 4.7 to gain 0.5. A consumer that only accumulates may still prefer shorter — $H=1$ returns 76% of the held-out benefit at under a quarter of the development cost, for 1.24 points of frame EPE — which makes the horizon a dial rather than an optimum.

I. Head-to-Head Against a Published Stabiliser

BiDASStabilizer [3] is the closest published plug-in stabiliser we could run. We compare on DRENDS, where neither method is in domain, on its own RAFT-Stereo *robust* prediction over all five recordings under one support neither influences (117.4M pixels). Two asymmetries run in opposite directions — the baseline is bidirectional, which our setting forbids; we are trained on surgical endoscopy, it is not — so we read the margin as an upper bound on our advantage.

Raw EPE 1.199 becomes 1.175 under the baseline and 1.131 under TETHER; Bad1 44.92% becomes 44.81 and 43.39, RMSE 1.876 becomes 1.776 and 1.710 — two refiners of a weak prediction (mean ground-truth disparity is 9.89 px), not a claim that either recovers a good one. TETHER is the only causal entry, at $62\times$ fewer trained parameters. The exception is named: on Vid13 TETHER is behind the baseline and 0.013 px *worse than the raw prediction it refines*, the one recording of five where the intervention is a net harm.

The temporal measure reverses the ordering, and we report it for that reason: change error falls 64.9% for the bidirectional baseline against our 42.5% at lag 1 (37.2 against 13.8% at lag 4), in every recording at every lag. A method that sees the future is more temporally stable and simultaneously less accurate on every geometric metric here: a reader given only a consistency score would rank these two the wrong way round.

TC-Stereo [6], run as its authors specify with its rigid pose warp on SCARED-C’s corrected poses, is outside its domain before any temporal component runs (frame path EPE 0.524 against 0.303 raw on the same support), and enabling the temporal path makes both metrics worse — a co-designed architecture meeting a scene its rigid warp does not describe rather than a weak method. Protocol, self-check and per-sequence numbers are in the supplement.

J. Ablating Our Own Design Decisions

Four design choices were ablated under a pre-registration naming held-out Bad1 as the endpoint and a three-seed spread of 0.37 points as the threshold below which a

difference is not read — the standard deviation over the base configuration’s three seeds of its pooled Bad1 reduction across all 35 backbone–sequence cells, computed before the registration was written; A4 remains a single run. A separate later registration governs *promotion* of a variant to the shipped model and decides that on development only, which is how the 38-channel head was promoted; the two rules are distinct and we apply each where it belongs.

TABLE IV

HELD-OUT RELATIVE REDUCTION (%), NOT ERROR: HIGHER IS BETTER, MACRO OVER BACKBONE–SEQUENCE CELLS. EACH INDENTED DEVIATION IS MEASURED AGAINST THE UN-INDENTED BASE ABOVE IT, AS ITS PRE-REGISTRATION NAMES. DROPPING THE LEARNED EVIDENCE FROM THE 142-CHANNEL BASE IS WHAT TETHER NOW IS, SO THE SHIPPED MODEL HEADS THE TABLE; ITS OWN FULL-RESOLUTION AND UNCONSTRAINED-OUTPUT DEVIATIONS ARE IN THE TEXT. † MARKS SINGLE RUNS.

Configuration	Evidence	D7 EPE red.	D7 Bad1 red.
TETHER (3 seeds)	38 ch	5.61 ± 0.42	5.84 ± 0.22
geometry only†	13 ch	4.44	4.64
+ learned evidence (3 seeds)	142 ch	5.06 ± 0.44	5.15 ± 0.34
no appearance†	78 ch	4.93	4.88
full-res context (3 seeds)	142 ch	5.24 ± 0.52	5.68 ± 0.16
unconstrained output (3 seeds)	142 ch	1.70 ± 1.43	6.97 ± 0.83

The result (Table IV) is not the one we expected, and it is why the shipped model is the smaller one. We began where the video-stereo literature does, with a frozen ResNet-18 [10] supplying appearance correlations and cost curves on top of the geometric channels — the configuration the four deviations were pre-registered against. Three sit above it on the held-out endpoint and the fourth, dropping appearance, sits inside the threshold: against the 142-channel base, none of the four registered decisions is load-bearing. Against the head we ship, two of the same decisions *are*: the quarter-resolution context loses 0.76 points when moved to full resolution (5.08 against 5.84 ± 0.22, one seed), and dropping the 25 correlations loses 1.20 (4.64, one seed). The learned evidence was absorbing the difference — with it present the architecture is insensitive to these choices, and removing it is what makes them matter. Removing it entirely, under the promotion rule that decides on D2 with D7 reported afterwards, stays ahead on both endpoints with seed deviation eight times tighter (0.32 against 2.53 points of D2 Bad1), and out of domain the ordering holds (5.20 ± 0.13% against 4.83 ± 0.15% DRENDS EPE, three seeds each on three backbones). The learned evidence is not what makes the module work; it is what makes its training unstable, at 157,504 frozen parameters, three forward passes per frame and 13.8% of the runtime. Its apparent advantage on the development closure (4.05% against 3.69%) is a sharp seed-0 artefact: seed 0 is the *best* of three for it and the *worst* for us, so the single-seed pairing anyone would report is the worst available to us, and over three seeds the ordering reverses (4.13 ± 0.38% against 3.55 ± 0.45%).

Two rows need the registered rule read out loud. Full-resolution context wins the registered endpoint against the

configuration it was registered against (5.68 ± 0.16 against 5.15 ± 0.34), an outcome the registration did not enumerate. On the base we ship, the same deviation loses 0.76 points; the registered configuration is also the better one.

Ablating convexity (A4) posts the largest held-out Bad1 reduction in the table, 6.97 ± 0.83%, but its margin over us (1.13 points) is not separable from its own three-seed spread — four times ours, and twice the registered threshold, which assumes comparable spreads and cannot be applied here. Its EPE loss *is* separable: 1.70 ± 1.43% against our 5.61 ± 0.42, over four standard errors. We decline it on the metric where the difference is real rather than on the metric that suits us. The two trajectories put side by side say what the constraint buys: from development to held-out, A4’s EPE gain falls 5.32 → 1.70 ± 1.43 while ours rises 4.13 → 5.61 ± 0.42, and the drop holds in all fifteen backbone–seed pairs — the unconstrained output buys in-domain accuracy that does not survive the split shift, and its 38-channel twin repeats the pattern (supplement). Its registered consequence we honour: the falsification clause required that if A4 matched or beat the canonical model held out, the claim attributing *cross-backbone* transfer to the convex action space be withdrawn rather than softened. It is withdrawn and appears nowhere in this paper; the *cross-split* statement above is a narrower claim, measured on A4’s own three seeds.

K. The Support Contract Is Load-Bearing

Confining the output to (1), not merely computing it, is a single line. On DRENDS the difference is categorical: unmasked, the shipped head *regresses* depth RMSE in every one of the 25 recording–backbone cells, by +221.9% on average (range +29.5 to +594.5); confined, it improves in every one of the same 25 cells, by −3.5% (−6.7 to −1.8). Zero-padded out-of-bounds memory enters the blend as $\tilde{m}_t = 0$ and bounded recurrence carries the collapse to the next re-anchor; the offenders are rare and bursty (300 of 38.8M pixels on one recording, aligned memory exactly 0.0 for every one) and, since invalid predictions on valid support are penalised rather than discarded, they are the entire source of a depth tail we had attributed to the external domain. On SCARED-C the same ablation is nearly inert: the contract costs a little in-domain gain, buys well-posed behaviour out of it, and is almost invisible to anyone validating in domain only.

VI. ANALYSIS

Alignment is necessary but not sufficient: the aligned memory carries a 12.74% oracle ceiling that indiscriminate use does not collect (Table II). The mechanism is alignment plus learned graded intervention plus bounded recurrence, not smoothing; the head never solves stereo matching, it chooses between two candidates.

a) *What each branch of $w_t = r_t f_t$ does.*: Suppressing either factor at inference, the trained head unchanged, separates their jobs. Fusion alone keeps most of the gain, 3.49% against 3.69%, but breaks a third more pixels (0.084% New-Bad1 against 0.063%); reset alone keeps 2.23% and breaks




Fig. 2. Where the module intervenes, and whether it helped there. Grey and off-white are off support, where (2) makes the module the exact identity and δe is zero by construction, so those pixels are not drawn at the centre of the diverging scale. **Top:** the median frame of a D2 sequence — acts on 19.2% of pixels at mean weight 0.46, improves 82.8% of them. **Bottom:** the worst frame of the same sequence, by rank rather than choice — 2.2% at 0.79, improves 12.9%; note the δe scales differ 40 $\times$ between rows (± 0.03 against ± 1.11 px). “Acts on” is the fraction of the frame’s pixels on support. Few pixels, high weight, wrong is a tail and not a trend: over the sequence’s 1,861 frames with defined weight, per-frame mean weight correlates *negatively* with harm (Spearman -0.68 ; -0.26 , same sign, on a held-out sequence).

0.211%. The fusion weight sets how much is gained, the reset gate how little is broken, and a paper reporting only the product and only mean error would have seen neither. Sec. V-H reaches the same division from accumulated geometry. Ablating convexity improves RMSE and worst case while collapsing the mean EPE gain (Sec. V-J): the failure is broad degradation of pixels already close, so the constraint buys mean accuracy rather than safety from large corrections.

Two limits are worth stating. Against the raw-versus-memory oracle of Table II — measured on TETHER’s own aligned memory, since under $H=4$ that memory is the previous fused output on three frames of four — TETHER recovers 29.0% of the available gain on D2 pooled and 25.2% per cell, so the open problem is deciding *where* to intervene.

And the effective weight w_t is an intervention magnitude rather than an uncertainty. Per pixel it ranks harmful intervention at AUROC 0.393–0.408 on D2 and 0.488–0.507 on held-out D7. The development figure is not chance-level but *inverted* — read in the benign direction it is 0.592–0.607, so a larger weight marks a pixel more likely to have been helped — while on held-out data the same reading gives 0.493–0.512 and orders nothing. The signal exists only on the split the checkpoint was selected on, which is the selection optimism of Sec. V-D appearing a second time, and does not survive to data the head was not tuned against. The 142-channel configuration has the same shape, so this belongs to the formulation and not the evidence space. It does not contradict the intervention ratios, which factorise into a count and a magnitude term needing no per-pixel ranking: a policy can be good on average and uninformative per pixel, which makes the oracle gap a ranking problem rather than a capacity one.

VII. DISCUSSION AND LIMITATIONS

a) What the splits support.: The closure is development work on D2: it supports mechanism analysis and the pre-frozen $H = 4$ operating point, not confirmation. One external domain does not establish universal robustness, DRENDS supplies a temporally filtered ToF reference rather than independent stereo ground truth, and per-backbone D7 margins come from one seed — the three-seed study bounds only the pooled effect.

b) What the metrics cannot say.: No-reference temporal consistency is gameable, shown on D4D [21], its only role here: the module reduces that score by 53.9%, copy-previous by 97.3%, warped-previous by construction to zero. The accumulated-geometry measure of Sec. V-H survives the same controls, which is why the robotics claim rests on it. Pull-warp support detects out-of-frame sampling but not internal disocclusion.

c) Where the intervention costs a map.: On development D2 the mapping measurement runs the other way: refinement *increases* world-frame disagreement by 15.5% at the shipped horizon, uniformly, all 15 cells negative. The account in Sec. V-H is consistent with every cell we have but is not proven: the two splits’ excess ranges never overlap, so no backbone makes D2 as hard as D7, and within D2 the per-backbone ordering does not follow excess either. What the grid does establish is that the reversal is a property of the split and of the horizon, not of one estimator: a module refining a prediction already close to its reference can cost a map more than it returns.

d) Cost, and where it is.: The head is compact but the deployed system also runs frozen stereo and two SEA-RAFT passes: 28.90 ms overhead at 95.2 MB peak VRAM on an H100. The frozen estimator dominates, so the end-to-end rate is set by the backbone, 11.6 FPS on Fast-FoundationStereo (overhead 33.5%) to 1.4 on StereoAnywhere (4.0%); *no real-time claim is made*. Stage timing locates it: the two SEA-RAFT passes are 26.2 of the 28.9 ms, against 0.67 for the evidence tensor and 1.00 for the head; the tensor is that cheap because TETHER builds no learned features, where the 142-channel ablation spends 5.13 ms there. The reverse pass exists only to compute C_t^{FB} , and the trained head barely reads it: pinning the cue to a constant at inference costs the shipped checkpoint +0.08 points of held-out Bad1 (5.77 against that same seed’s 5.69, inside the seed spread), and a head trained *without* the cue is equally indifferent to receiving the real one (-0.02), so the indifference is architectural rather than an accident of one checkpoint. Training without it does cost accuracy (supplement), so the cue is a training-time scaffold: **train with it, deploy without the reverse pass**, taking the module from 28.9 to ~ 15.8 ms on the unchanged checkpoint. Both held-out arms are single seeds under a pre-declared threshold. We evaluate perception geometry and what a map accumulates from it, not downstream control, and demonstrate neither clinical safety nor patient benefit.

VIII. CONCLUSION

TETHER attaches one causal temporal module, trained once, to frozen and heterogeneous stereo estimators without touching their internals or using future frames. Its gain is not temporal averaging: matched fixed and EMA blends under an identical state machine recover a sixth of it on development — paying in bad pixels there — and a ninth out of domain. One frozen $H = 4$ checkpoint improves all five estimators on held-out D7 and all 25 recording-backbone cells zero-shot on DRENDS.

Five results bound what may be claimed, and they are why the protocol is part of the contribution. No-reference consistency is won outright by replaying warped history while the accumulated-geometry measure is not, so the robotics claim rests on the second. The fusion weight ranks harm only on the split it was selected on: a magnitude, not an uncertainty. Selection optimism belongs to the evidence space rather than the protocol. A module that declares a support must apply it to what it emits. And what a map fuses improves held out while degrading on development, the term responsible being accumulated alignment residual — so the horizon is an exchange rate, not an optimum. What remains open is *where* to intervene.

REFERENCES

- [1] Z. Li *et al.*, “Temporally consistent online depth estimation in dynamic scenes,” in *Proc. WACV*, 2023.
- [2] N. Karaev *et al.*, “DynamicStereo: Consistent dynamic depth from stereo videos,” in *Proc. CVPR*, 2023.
- [3] J. Jing *et al.*, “Match Stereo Videos via Bidirectional Alignment,” in *Proc. ECCV*, 2024.
- [4] Y. Wang *et al.*, “PPMStereo: Pick-and-play memory construction for consistent dynamic stereo matching,” in *NeurIPS*, 2025.
- [5] Y. Wang *et al.*, “SEA-RAFT: Simple, efficient, accurate RAFT for optical flow,” in *Proc. ECCV*, 2024.
- [6] J. Zeng *et al.*, “Temporally consistent stereo matching,” in *Proc. ECCV*, 2024.
- [7] H. Xu *et al.*, “StereoDiffusion: Temporally consistent stereo depth estimation with diffusion models,” in *Proc. MICCAI*, 2024.
- [8] F. Aleotti *et al.*, “Neural disparity refinement for arbitrary resolution stereo,” in *Proc. 3DV*, 2021.
- [9] W.-S. Lai *et al.*, “Learning blind video temporal consistency,” in *Proc. ECCV*, 2018.
- [10] K. He *et al.*, “Deep residual learning for image recognition,” in *Proc. CVPR*, 2016.
- [11] L. Lipson *et al.*, “RAFT-Stereo: Multilevel recurrent field transforms for stereo matching,” in *Proc. 3DV*, 2021.
- [12] J. Li *et al.*, “Practical stereo matching via cascaded recurrent network with adaptive correlation,” in *Proc. CVPR*, 2022.
- [13] L. Bartolomei *et al.*, “Stereo Anywhere: Robust zero-shot deep stereo matching even where either stereo or mono fail,” in *Proc. CVPR*, 2025.
- [14] J. Min *et al.*, “S²M²: Scalable stereo matching model for reliable depth estimation,” in *Proc. ICCV*, 2025.
- [15] B. Wen *et al.*, “Fast-FoundationStereo: Real-time zero-shot stereo matching,” in *Proc. CVPR*, 2026.
- [16] Y. Wang *et al.*, “NVDS+: Towards efficient and versatile neural stabilizer for video depth estimation,” 2023, arXiv:2307.08695.
- [17] Y. Geifman and R. El-Yaniv, “SelectiveNet: A deep neural network with an integrated reject option,” in *Proc. ICML*, 2019.
- [18] M. Allan *et al.*, “Stereo correspondence and reconstruction of endoscopic data challenge,” 2021, arXiv:2101.01133.
- [19] P. J. Edwards *et al.*, “SERV-CT: A disparity dataset from cone-beam CT for endoscopic 3D reconstruction,” *Med. Image Anal.*, vol. 76, 2022.
- [20] Y. Long *et al.*, “E-DSSR: Efficient dynamic surgical scene reconstruction with transformer-based stereoscopic depth perception,” 2021, arXiv:2107.00229.
- [21] R. Docea *et al.*, “The Dresden Dataset for 4D reconstruction of non-rigid abdominal surgical scenes,” 2026, arXiv:2603.02985.
- [22] G. Loza *et al.*, “DRENDS: A dataset for depth in robotic endoscopy with dynamic scenarios,” Zenodo Dataset, 2025, doi: 10.5281/zenodo.17598453.
- [23] Z. Teed and J. Deng, “RAFT: Recurrent all-pairs field transforms for optical flow,” in *Proc. ECCV*, 2020.
- [24] W. Hu *et al.*, “DepthCrafter: Generating consistent long depth sequences for open-world videos,” in *Proc. CVPR*, 2025.
- [25] M. Poggi *et al.*, “Continual adaptation for deep stereo,” *IEEE TPAMI*, vol. 44, no. 9, 2022.
- [26] M. Poggi *et al.*, “On the uncertainty of self-supervised monocular depth estimation,” in *Proc. CVPR*, 2020.